

Taxis in the Mist: Predicting Ride Demand

John Blakkan, Rohit Nair, Andrea Pope

December 5, 2016

W205Project

UCB W205 Group Project by John Blakkan, Rohit Nair, Andrea Pope



Figure 1: Boroughs of New York City

Purpose and use-cases

Our goal is to model demand for taxi and ride service vehicles in New York City. We are not concerned with actual cab dispatch, but want to predict overall borough demand (on an hourly basis) so that cab companies and ride hailing services can pre-position cabs (and cab service staff) at appropriate depots.

We model cab (and ride-service) demand as a function of week-of-year, day-of-week, hour-of-day, and weather factors.

Architecture

There are seven main pieces. Refer to figure 2.

1 - Ride data ETL

This is consolidated for several taxi/ride-hailing services (in hive). Fundamental fields are: Date/Time of pickup, Lat/Long location of pickup, and number of passengers. For our purposes, fare information, taxes, etc. are not retained)

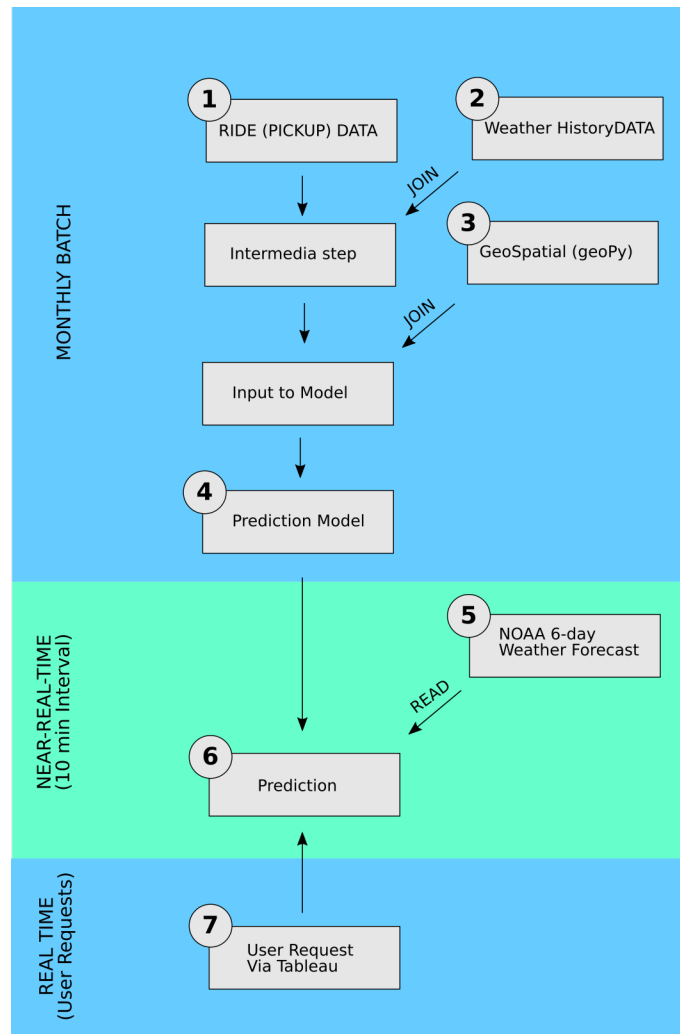


Figure 2: Architecture figure

2 - Historical Weather ETL and join to the Ride Data

This is joined to the ride data, giving historical ride data with weather (in hive, as measured at NY LaGuardia Airport). We get this data from NOAA at <https://www.ncdc.noaa.gov/cdo-web/>. Since we use this only for model building (a batch process), we only need to pull it when the model is updated (e.g. not when current weather is updated). NOAA requires this data be accessed by a request for data return by email.

The data (ingested in .csv form) is lightly transformed (in excel or equivalent, although this could be scripted.) Snowfall is accounted as rainfall in a 10:1 ratio (i.e. 10mm of snow = 1mm of rain). Fields used are: Date, Total Precipitation (in tenths of mm), Daily Max temp (F), and Daily Min temp (F)

We implement a hash-join to join weather information to the ride data with pyspark code because the daily weather data is very small compared to the ride pickup table (i.e. there are many taxi pickups per day, but we only have one overall weather report (with 24-hr max/min temp and precipitation) per day. In-day variances are co-linear with hour-of-day, so will have input to our model.

Pyspark map/lambda is used to perform this hash-join.

3 - Local Region Identification transformation of Ride data

Our customer use case is to plan allocation of taxi/ride-service vehicles to smaller regions of NYC, such as zip-codes or boroughs, so the lat/long co-ordinates in the Ride Data must be converted.

One method we used is to do this in pyspark, with a program called `hive_borough.py`. In turn, it runs `borough_finder.py` on each row of the Ride data. `Borough_finder.py` (including its table of borough-bounding-polygons) is distributed by pyspark to all potential parallel processes (much like broadcast variables). See the test script for example use (including counts using pyspark accumulator variables).

We obtain this data by using the online SAAS at <https://www.itouchmap.com/latlong.html> to create bounding polygons for each borough. Boroughs can be bounded with approximately 20 vertices (with water borders requiring fewer points, and complex street boundaries - such as between Brooklyn and Queens requiring more). Note: The points are defined by the project, not by the website, so there are not IP or Terms-of-Use issues. The New York City government website was used as a reference for the borough borders (<http://maps.nyc.gov/doitt/nycitymap/>).

Here is a sample of one of the bounding polygons, for the “pseudo-borough” of Eastern New Jersey. (Which is obviously not part of New York City, but our analysis showed significant NYC taxi/ride pickups here, so we include it. (See comment in the section “Look at the raw data”, below)

These immutable polygon vertices can be used with matplotlib’s point-in-polygon library to provide a transform function for pyspark’s map-lambda to add borough identifiers to the Ride data.

Later in the project we determined our preferred method, however. We currently use the geoPy module for reverse geocoding. We were initially concerned that, because this module delegates reverse geocoding to a variety of online services, that it would be too slow for our volume of data. However, the module’s internal caching of regional co-ordinates ultimately proved to have acceptable performance, so is used for it’s zipcode level resolution.

4 - Predictive model, with parameters

This is produced from the ride/weather data, and placed in a (TBD) table (in postgres).

5 - Current weather forecast for NYC

This will be pulled every 10 min from NOAA’s XML service (http://graphical.weather.gov/xml/sample_products/browser_interface/ndfdBrowserClientByDay.phpusing). Script `forecast_server.py` parses the XML




Figure 3: Bounding Polygon for Eastern New Jersey region

and creates a postgres table 6-day `weather_prediction`. `forecast_server.py` can be run from the command line, but for production will be run via `start_forecast.py`. This script cleans the data by translating the mixed 12hr/24hr results into uniform 24 hr. results, and transforming pairs of 12 hr precipitation probabilities into single 24hr probabilities (i.e. the 24 hr. probability is the max of the day and night probabilities). Timestamps of the NOAD downloads are also kept, so that “stale” weather data can be identified. We pull and cache (and transform) the data in this manner because we have observed that the site has a high variance in response time (up to 10s of seconds) and also becomes unavailable for minutes at a time. This is only a 6 line table (“`weather_prediction`”) so we use postgres. Attention is paid to using transactions when modifying this table.

6 - Prediction

Then script `ride_prediction.py` takes the forecast from postgres table “`weather_prediction`” and the model parameters from postgres (TBD) table, and produces postgres table `ride_prediction`. `ride_prediction.py` can be run from the command line, but for production will be run via crontab (at 5, 15, 25, 35, 45, and 55 min after the hour). The output of this table is: For each of 5 boroughs, for each of 6 days, for each hour in each day, a predicted number of large taxis (4 or more passengers) and small vehicles (3 or fewer passengers). The timestamp of the `weather_forecast` used to generate the prediction is also passed through, so the user can verify the age of the forecast used for the predictions.

This 720 row table (5 Boroughs * 6 Day Forecast * 24 hrs/day) is kept in postgres.

7 - Tableau presentation

Tableau will serve data from the `ride_prediction` table. The predicted number of pickups

- By zipcode
- For each hour
- For the next 6 days

will be presented (along with appropriate graphs).

Users will also see the (UTC) timestamp of the last successful fetch of weather forecasts, to be informed in the event that the forecast is more than a few minutes out of date.

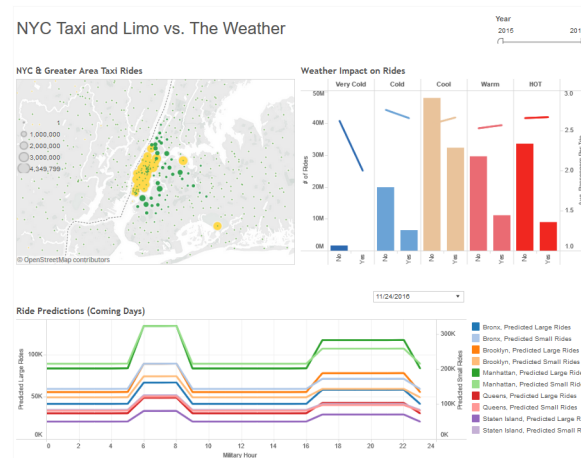


Figure 4: Tableau Screenshot

Dependencies and operation

The standard EC2 AMI is used, running on an M3 Large instance.

Python modules

- matplotlib - used for the “point-in-polygon” determination of boroughs from lat/long co-ordinates
- psycopg2 - postgresql access
- urllib - fetching the real-time weather forecasts from the NOAA XML server
- geopy - reverse geocoding (in some versions of the application, and for research/analysis)
- lxml - XML parsing of the NOAA real-time weather forecast.

Also used are standard modules, including sys, random, numpy, time, datetime, and threading (used for the weather forecast periodic fetch)

One time setup processes

Determine bounding polygons in lat/long co-ordinate system for the five boroughs of New York City (plus a region of Eastern New Jersey) was a one-time only setup. Ultimately we used a different method to perform reverse geomapping - the geoPy module, but we note this here as a one-time step, should it ever be necessary to revert to the bounding polygon method.

Monthly Batch process

This may be run (as user w205) via script XXXXX-TBA-XXXXX.py. The ETL flow is:

- 1) Load monthly ride data from source .csv files on Amazon S3. We're currently using July 1, 2014 to June 30, 2016 as our model training period. (If users choose to retain the raw .csv files, each month only one additional month's data will need to be fetched).
- 2) Strip header lines from ride .csv files
- 3) Copy into single hive table with consolidated, uniform format for all ride services
- 4) Load historical weather from source (???) .csv files
- 5) Strip header lines from weather .csv file
- 6) Hive join of weather onto ride data, based on date, creating table: XXXXX_XXX_XXX
- 7) Use borough_finder pyspark program to (spark-submit XXXXX.py) to convert lat/long information in XXXXX_XXX_XXX into borough names in hive table YYYY_YYYY

Periodic results pre-calculation

To speed user access, we pre-cache (every 10 minutes) the results of a 6-day weather forecast from NOAA, parsing it and copying it into postgres table weather_prediction. This is done from start_forecast.py. (This is a threaded script, run in background. An alternative is to run it as a cron job, but this wrapper script was found to have simpler administration). Note that this is done both for performance and reliability, as the NOAA site, in our experience, sometimes is unavailable for several minutes at a time. With the local copy, results are always available, even if slightly out of date. (Timestamps are available for the user to be aware if they are seeing "stale" weather forecasts. In practice, we've never seen a forecast go stale for longer than a single 10 minute interval)

In addition to the periodic update of the weather forecast, the ride-utilization model is run every 10 min, providing pre-computed results and eliminating user wait time. The scope of our problem permits us to pre-calculate for all 5 boroughs plus Eastern New Jersey.

Serving results

Results are presented via Tableau from the pre-calculation

Analysis, Observations, and Ideas for further development

Notes on Data

We believed our dataset covered New York City originating taxi rides. However, mid-way through the project, we discovered there were many rides originating outside New York. In particular, New Jersey had more "New York City" rides that did the borough of Staten Island. This caused us to adjust the project to also cover an Eastern portion of New Jersey, treating it (for purposes of analysis) as a "pseudo-borough" of New York city.

We also found other outliers, including rides originating in Las Vegas. These may include use of ride hailing services where the ride was actually booked while in Las Vegas, but we did not pursue further analysis of these, choosing simply to exclude them.

Notes on Data Transformation

We followed two parallel paths in determining how to translate the ride pickup location (in decimal lat/long co-ordinates) into local region. Our first method involved using public mapping sites to produce bounding polygons for the boroughs. This involved a one-time effort, and required no web access during later processing. Library algorithms to do "point-in-polygon" determination were used. We also added a speed enhancement to this: Each datapoint was roughly categorized as "Western," "Northern," or "Southern," and this initial categorization was used to determine the order in which the detailed polygon match was attempted. (For

example, if a point was Northern, the Bronx and Manhattan polygons would be checked first. If a point was Western, NJ and Staten Island would be checked first).

The second was to use available python modules (including geoPy) which could do reverse lookup. This was a two stage process: First lat/long to zip code, then a translation of zip-code to borough and neighborhood. The modules looked up data from public websites, but also did caching of geospatial information to maintain performance.

The first point-in-polygon was satisfactory with respect to performance, but we are currently using the geoPy/Zipcode method.

Future work and Extensions

Scale

If we were developing a nation-wide system, pre-calculating all the results could be problematic. Our scope was only 5 boroughs plus New Jersey. In the case of a later nation-wide expansion, we would likely include a different strategy - either distributing a full pre-calculation over many nodes, or perhaps by caching results for only the largest metro areas, with rural areas calculating only “on-demand” (with a correspondingly lower service-level-agreement on response time).

Resolution of weather to small regions

We regarded weather as uniform in our geographical region. The New York metro area is on the east coast of a continent, and generally has the a uniform weather pattern. This is on contrast to other areas, such as the San Francisco Bay, where there are significant regional microclimates. Expansion of the system to other metro areas might need to account for regional microclimates.

Observations and notes on the development process

Data Checkpointing is important

When working with large datasets (e.g. up to billions of raw records), we had several instances of multi-hour Data Ingress processes fail, resulting in up to half-day development delays. This happened for several reasons: Once when users - not realizing the task was underway, shut down the EC2 instance. Other times when spark jobs were running with the same application name (which, if not unique in the system, will cause jobs to silently lock up).

“Mock” databases (or subsets of larger database) are important

Even the monthly partitions of the databases are large enough that some steps will require many minutes of processing during the development phase. It was useful to be able to reduce the size of the input .csv files to greatly reduce the time taken by development test runs. To facilitate this, we produced a small script called “decimate.py.” It takes an argument indicating what proportion of the .csv lines in a file to retain (we often used 0.01, for 1%). Because we didn’t want our sampling to possibly correlate with any data sequence feature, decimate.py sieves out lines of a .csv stochastically, rather than by stepping through the file with an interval. We also incorporated a random seed argument, to assure we could make the decimation repeatable if necessary.

Rapid prototyping and unit testing are important.

The near-real time element of the project is the fetching of NOAA forecast data for the prediction model. Early in the project, we mocked up a simple Ruby-on-Rails website to drive our “poll of NOAA and parse” functions. Because this tool made it so easy to repeatedly test the NOAA interface, we found multiple problems: (1) NOAA actually provides only 6.5 of forecast data for some weather parameters, not 7 days. We stopped applications from erroring-out by using only the actual valid data (we picked 6 days). See the figure below, note the missing value for predicted min. temperature on day 7.

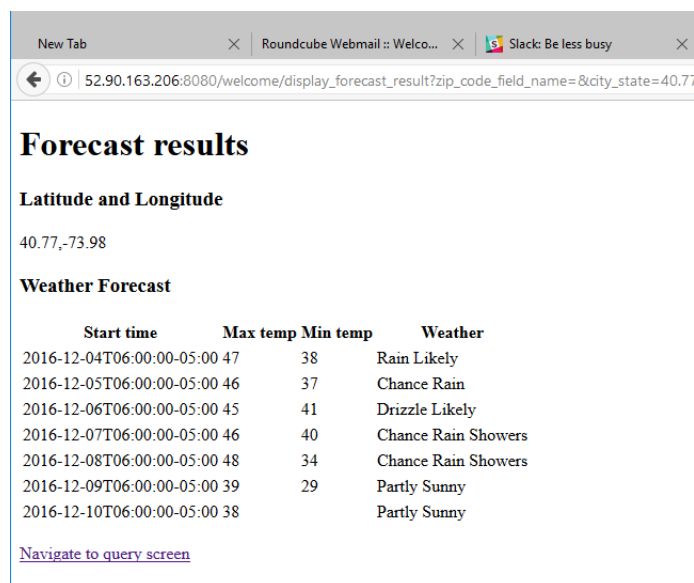


Figure 5: NOAA Weather Forecast: Ruby-on-Rails output

- (2) The NOAA interface has a highly variable response time - sometimes over 15 seconds. The NOAA site overall appears to go down for minutes at a time, particularly in the early hours of the morning (Eastern Time). [It is also possible that the NOAA website was recognizing too-frequent requests and throttling us; we never determined that.] Based on this, we realized we needed to cache a local copy of the forecast, and periodically refresh it. Had we not learned this early in the project, we would have developed a system which appeared to work in testing, but under higher load would have frequently failed.

Building the “Steel Thread” end-to-end preliminary version is a good idea. We didn’t do it.

The assignment recommends getting a stripped-down version of the full system (“A Steel Thread”) working early on, then adding to it. We did our integration relatively late, after building ETL and Serving systems. We did benefit from having mock databases to keep all elements being developed towards the same schema. But we did have an integration phase at the end of the project which was lengthier than it might have been if we’d integrated earlier.

Developing features in parallel with competing implementations is good (if resources are available)

One of the fundamental features of our project was the need for reverse geocoding. Our ride pickup data was all in lat/long co-ordinates. We needed to translate this into borough names. We *preferred* to translate it

down to zip-code resolution (i.e. neighborhoods). We proceeded on a parallel track: One effort used a set of bounding polygons around each borough, then used in-memory calculation of point-in-polygon (using pyspark, and shared user defined functions shared between all worker instances - see `borough_finder.py` and `hive_borough_demo.py`). It was fast (120 million conversions and writes to table in 48 min {with perhaps a third of that time in extraneous debugging output}), but our preferred solution was to use the geoPy module to convert to zip-code resolution. geoPy caches locations, so ultimately the solution we used was geoPy.

Look at the raw data, there may be surprises.

A minor issue we found was the presence of a small number of “out-of-new-york-city” pickup locations. Some were in Las Vegas, others were scattered around the country. We think this may be from user error in some application (i.e. having the location where the trip was booked, rather than the actual pickup location).

The more profound thing we found: The number of pickups in all of Staten Island appears to be smaller than the number of “New York” pickups that actually originate in New Jersey. This finding led us to re-consider the scope of the project; we consider eastern New Jersey to now be in-scope.

Other interesting items in the data (when perusing with HUE and Tableau): Ridership is greater in the fall than the spring; and passenger loads per cap seem higher during particularly warm or particularly cold weather.

“Skew” of data size needs to be considered; in-memory hash-joins may be better than general database joins

An hourly prediction of a year’s worth of activity is only 8760 elements. Rather than doing a database join operation, pyspark (with broadcast variables of hashes, or hashes included in the shared functions) may be a better choice. This was implemented in the `hive_borough.py` script.